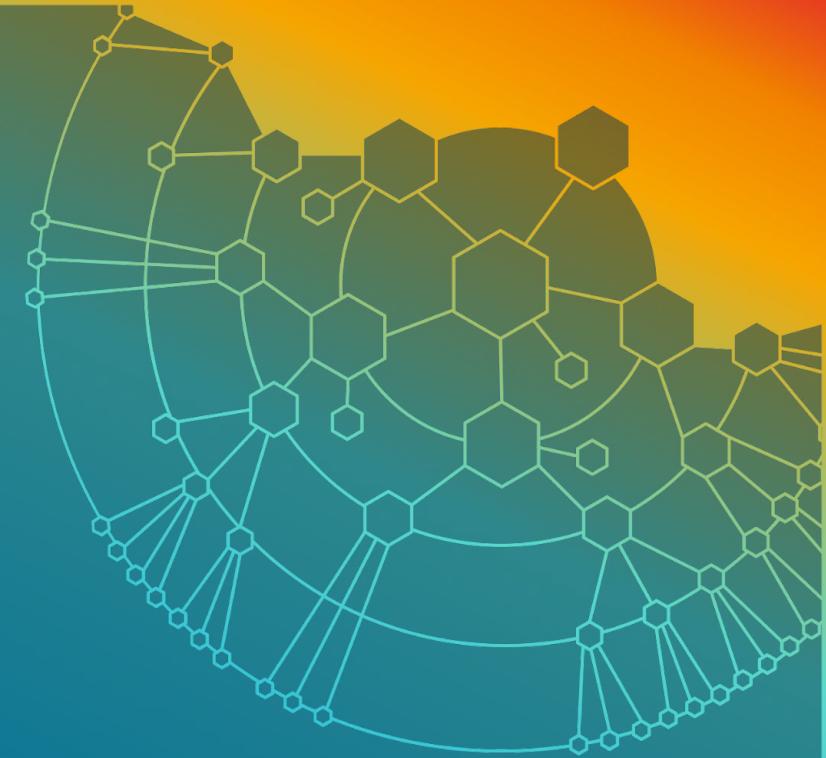


Thunderdome

Ian Davis, Protocol Labs



IPFS Camp

Thunderdome?



IPFS Camp

Thunderdome

A gladiatorial arena where ~~conflicts are resolved by a fight to the death~~ performance of different IPFS Gateway configurations and implementations can be compared



IPFS Camp

What does Thunderdome do?

- Runs experiments on demand and records metrics for the things being tested
- Currently it supports testing IPFS gateway implementations with HTTP traffic
- Designed to run for long periods - days
- Given an experiment definition it starts all the necessary infrastructure, applies load and monitors





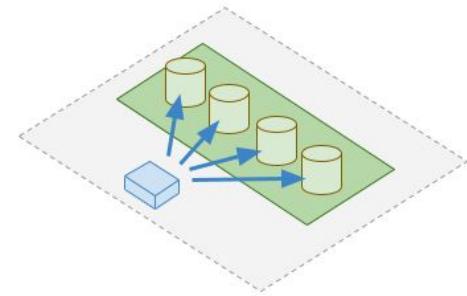
Thunderdome Infrastructure



IPFS Camp

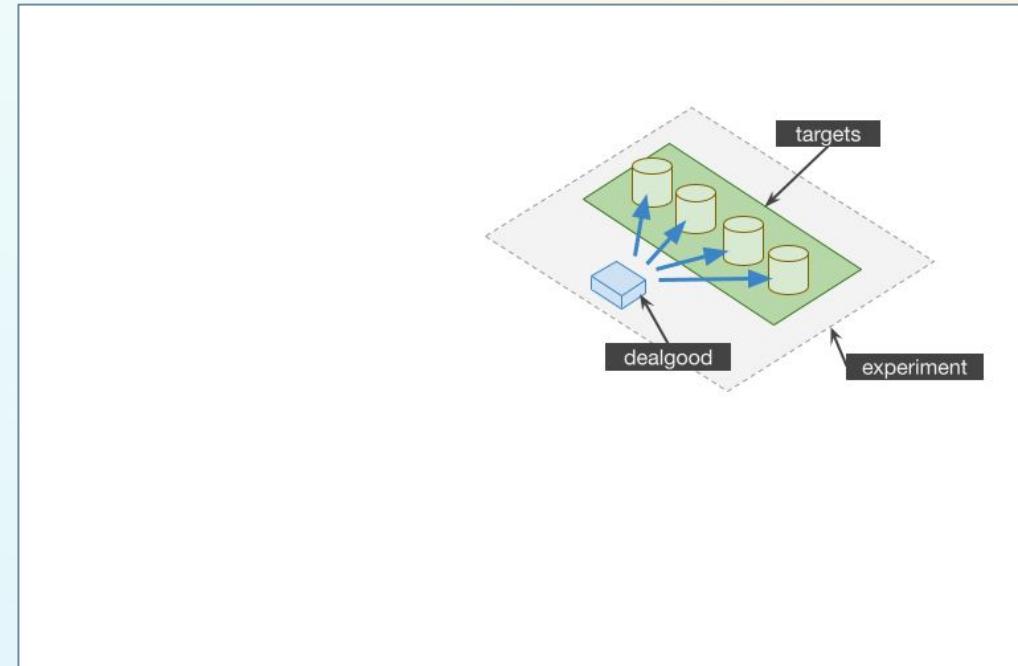
What is an experiment?

- At its simplest: fire requests at a bunch of ipfs nodes
- Today these are HTTP requests intended to test IPFS gateways



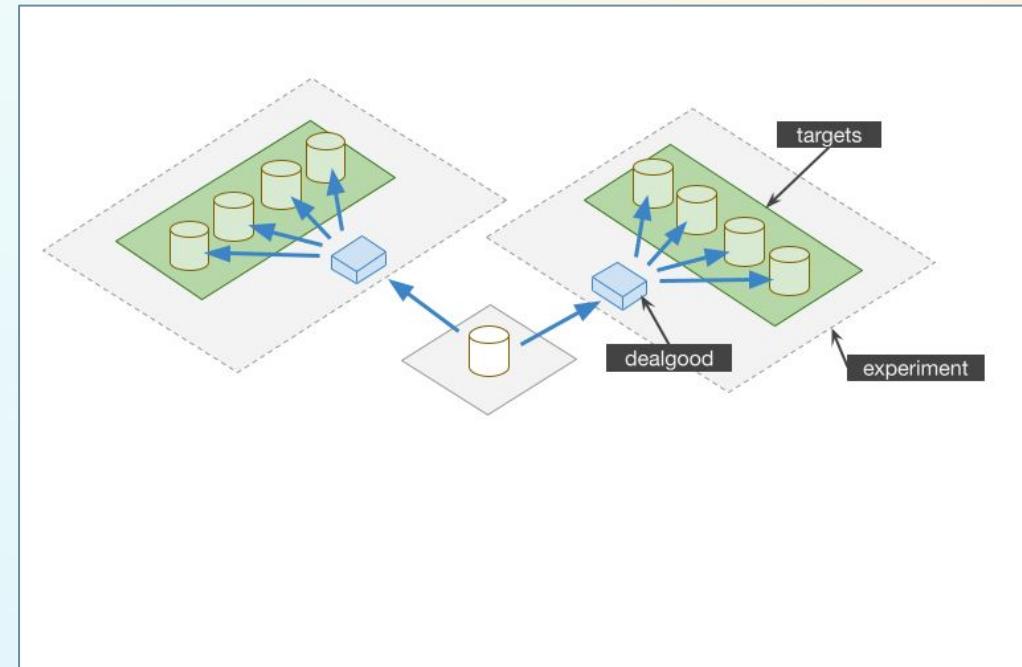
What goes into an experiment?

- **experiment** defines
 - request rate
 - concurrency
 - target config
- **target** - a docker container in ECS
- **dealgood** - applies load evenly to targets



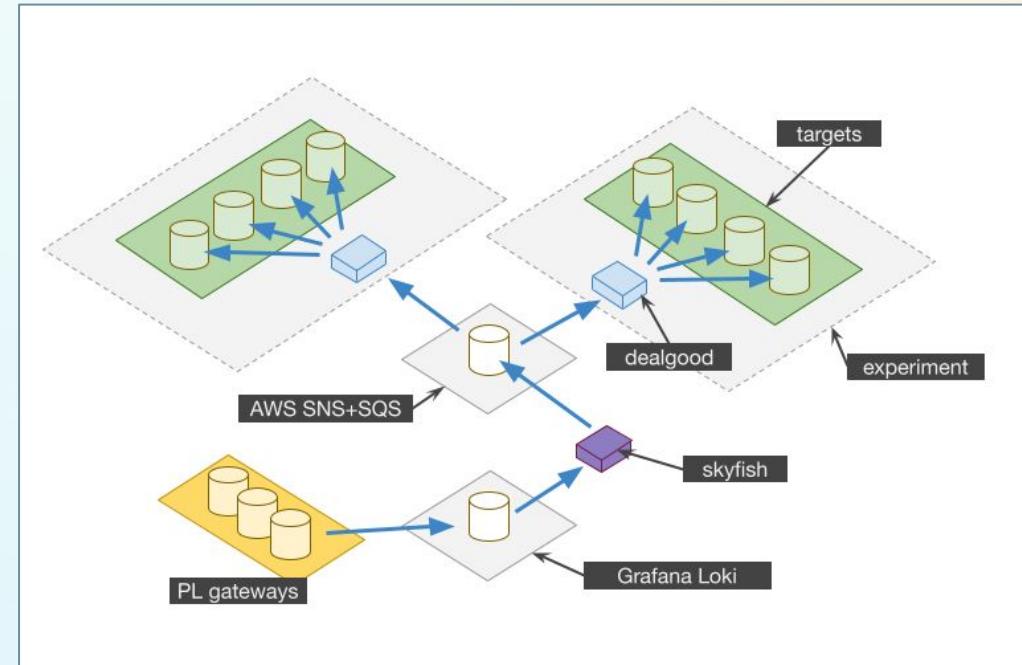
Multiple experiments

- Each experiment is independent
- Each instance of dealgood controls the load applied to that experiment's targets



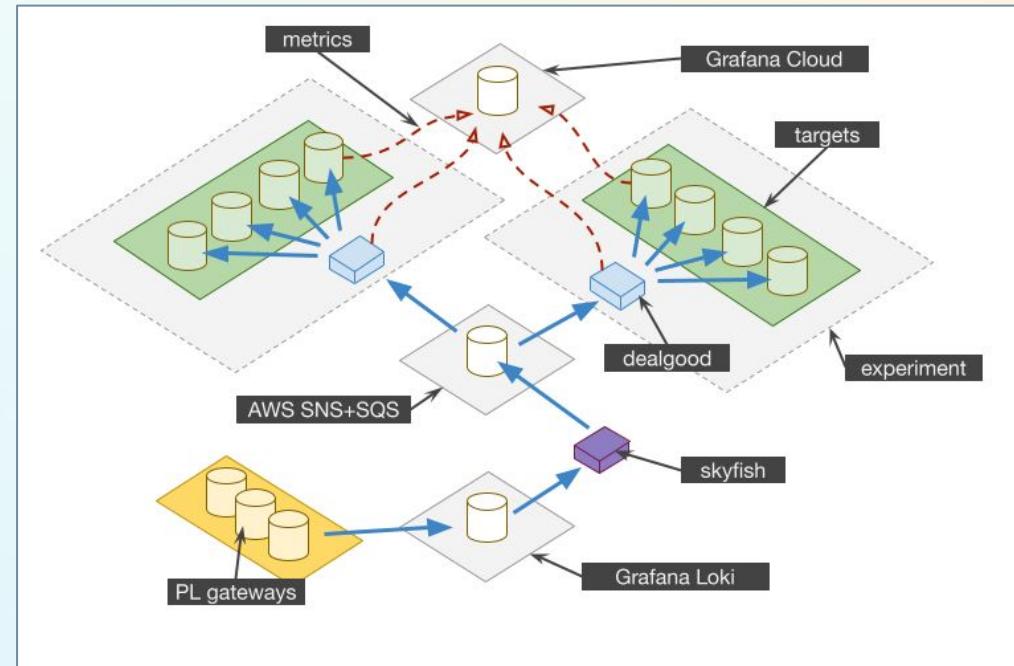
Where do requests come from?

- HTTP requests are sampled from ipfs.io gateways into Loki
- **skyfish** - bridges Loki and AWS SNS
- **dealgood** - listens to its own SQS queue fed by SNS



Metric collection

- Metrics are read from target containers and dealgood
- All available in Grafana cloud for dashboarding

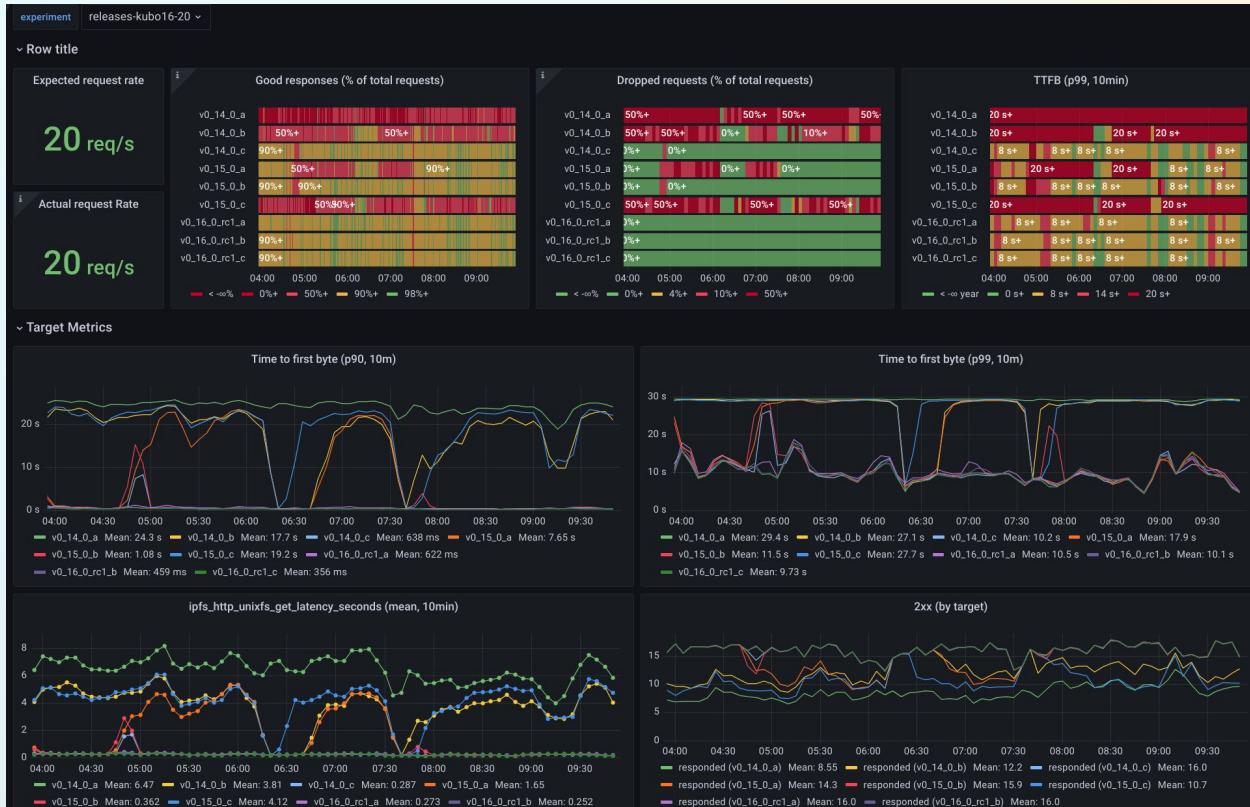


Experiment results

- Grafana holds all relevant metrics from an experiment
- Collected via prometheus from:
 - targets
 - dealgood
 - skyfish
- Tagged by experiment and source (job)

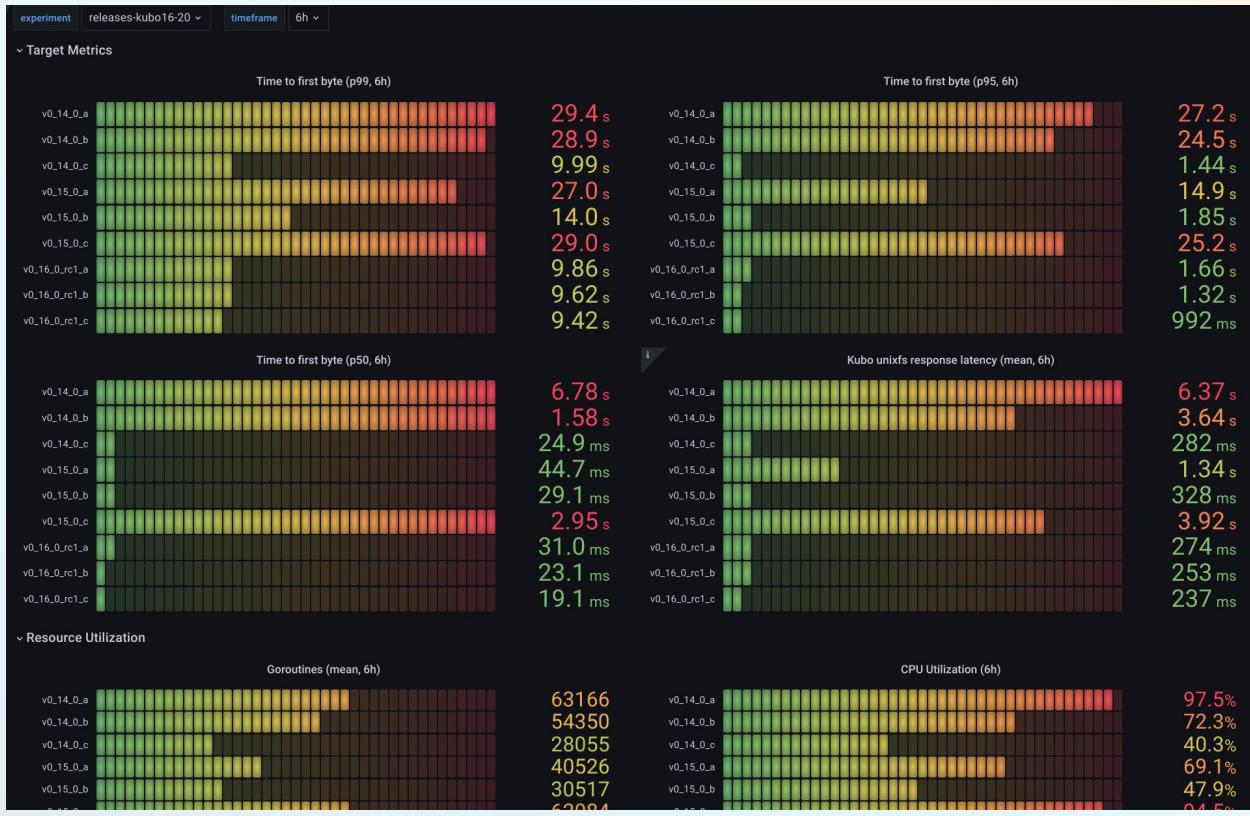


Experiment timelines



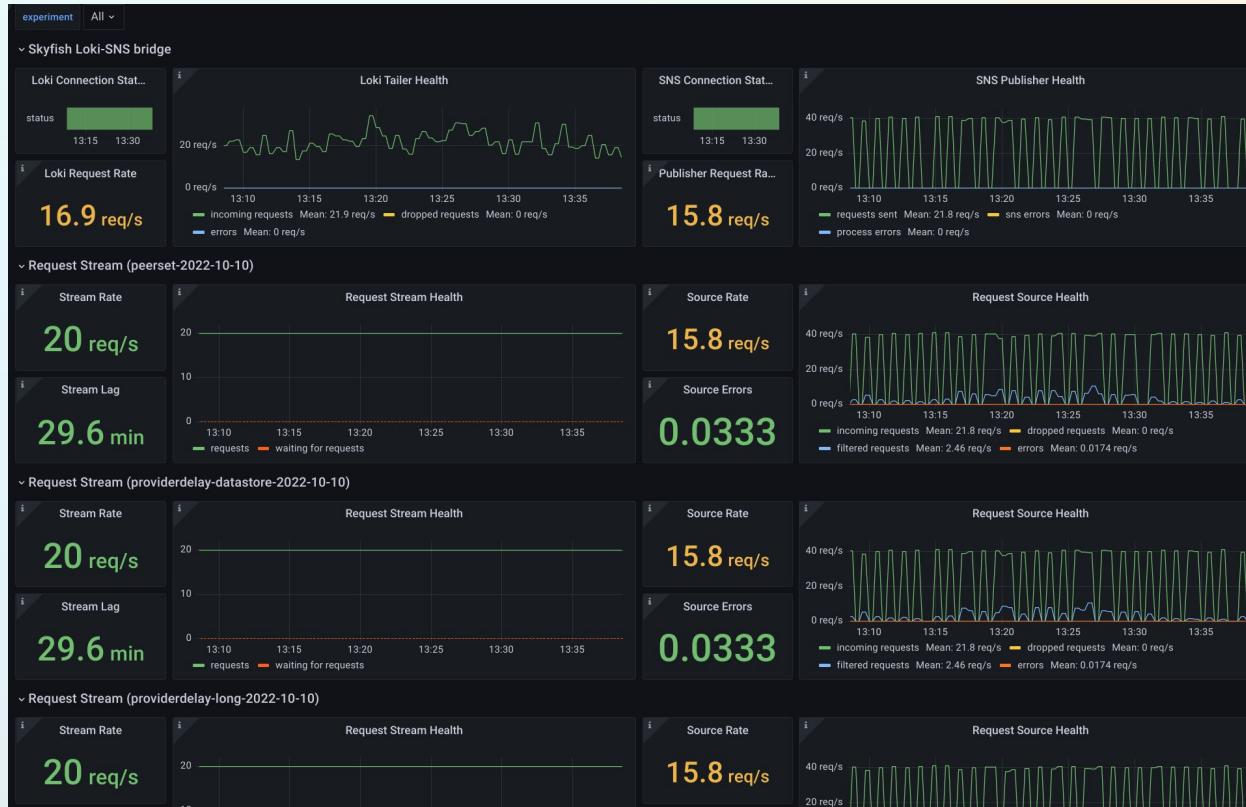
IPFS Camp

Experiment summaries



IPFS Camp

Experiment health



Infrastructure challenges

- The nature of ipfs and p2p makes it hard to test instances in isolation
- Nodes like to chat to their neighbors
- At a minimum we want to avoid test targets from fetching from one another
- Network ACLs isolate each target
- Tracking IPFS work on peer connection rules





Thunderdome Experiments



IPFS Camp

Initial experiments

- Our focus has been on proving Thunderdome
 - Ensuring no bias in infrastructure
 - Scaling request streams
- Picked some easy experiments:
 - Baseline experiment
 - Comparing kubo releases
 - Bitswap provider delays



Experiment: tweedles

- Baseline experiment
- Two identical Kubo instances
 - tweedle-dee and tweedle-dum
 - Same configuration
 - Same load
 - Should perform identically
- Detect bias in Thunderdome infrastructure



P99 TTFB



IPFS Camp

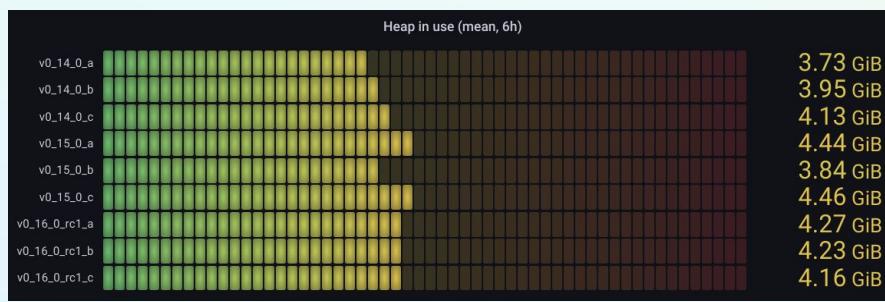
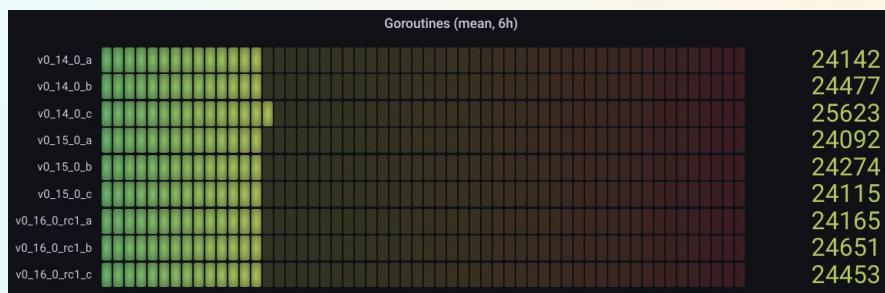
Experiment: kubo v0.16 release

- Version comparison
- Two pairs of each Kubo version:
 - v0.14, v0.15 and v0.16 rc1
 - Same configuration
 - 10 rps and 20 rps
- Detect potential performance regressions or improvements



Comparison at 10 rps

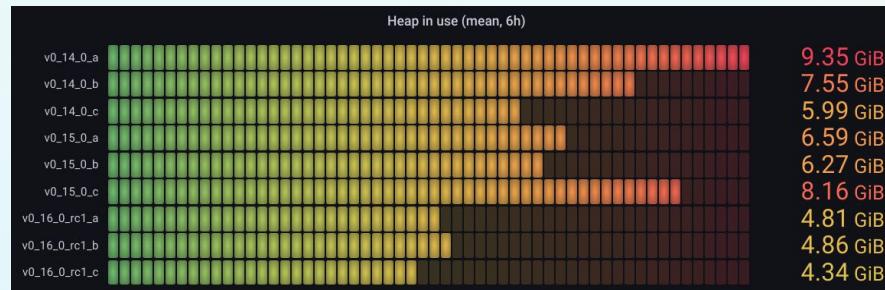
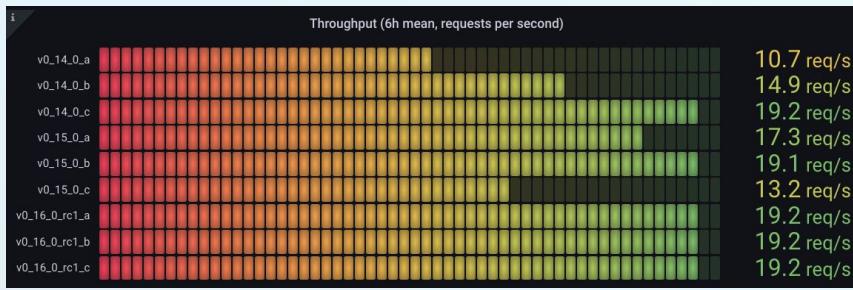
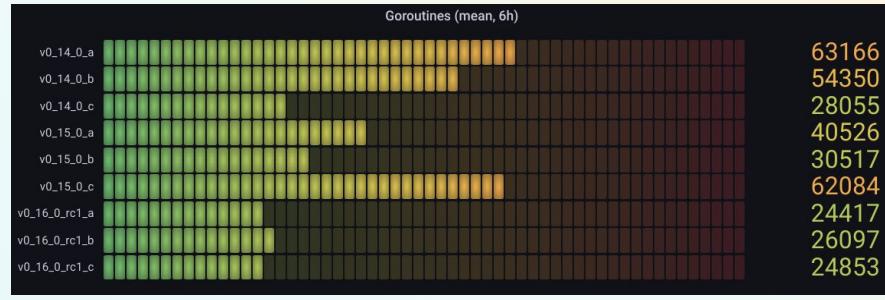
All quiet...



IPFS Camp

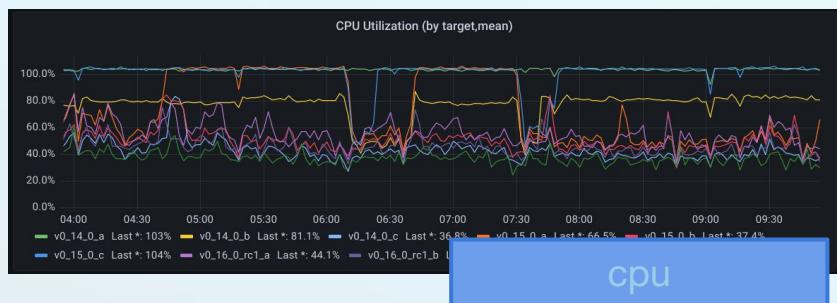
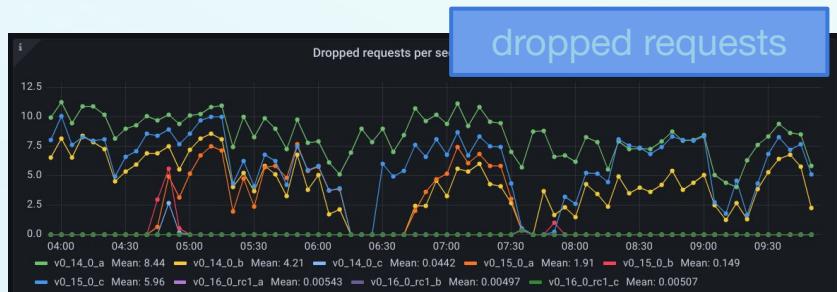
Comparison at 20 rps

Some things blew up...



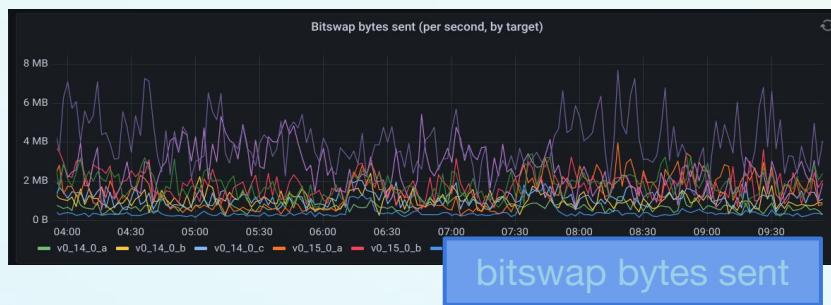
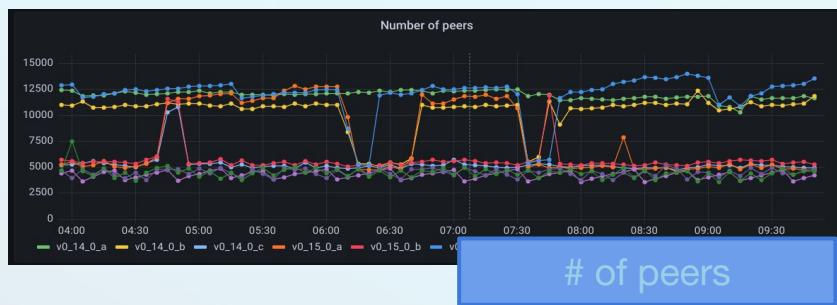
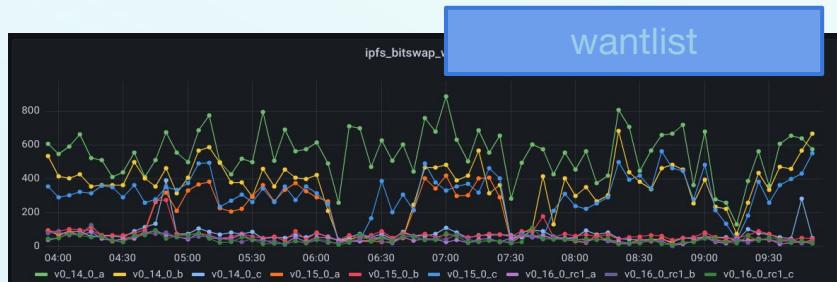
Comparison at 20 rps

Over time...



Comparison at 20 rps

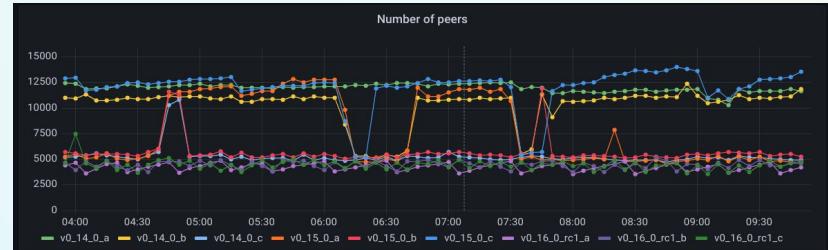
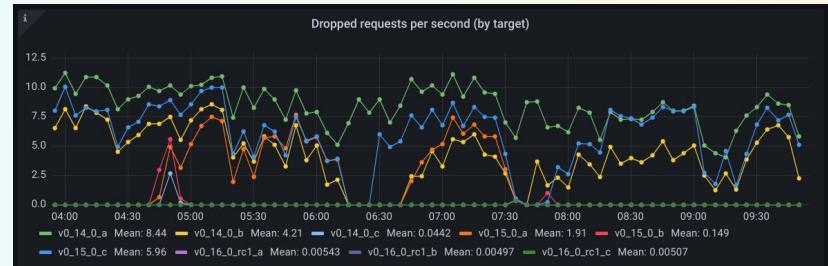
Bitswap behavior - elevated peer lists



IPFS Camp

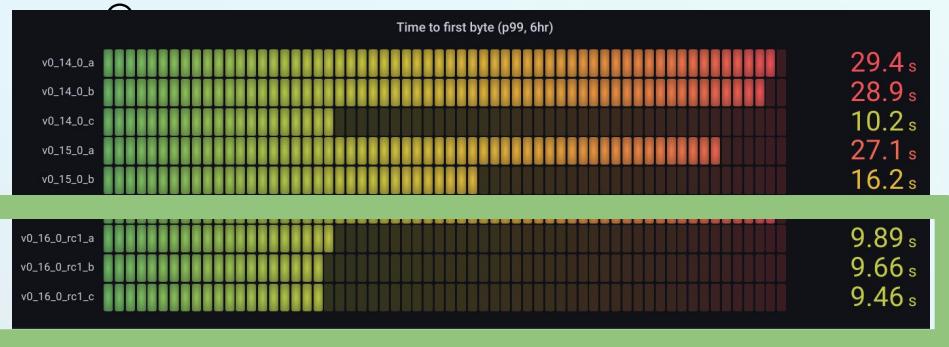
Observations

- Kubo can get trapped in a state with elevated numbers of peers which causes:
 - Large heap
 - High goroutines + CPU
 - Larger wantlists
 - High TTFB
 - Poor throughput



Conclusion

- But... Kubo 0.16.0 rc1 seemed to avoid falling into this state



Follow up

- Followed up with comparison of various commits between v0.15 and v0.16
- Narrowed down point at which Kubo appears to behave better - routing logic changes
- Read the experiment summaries
 - <https://tinyurl.com/45mhhf89>



Other Experiments

- **gwblock** — Compared performance of Kubo when instances are blocked from connecting to public gateway instances
- **providerdelay-base** — Compared various bitswap provider delay settings
- **providerdelay-peerset** — Compared various provider delays with high water of peer connections of 600
- **providerdelay-acceldht** — Same as base provider search delay experiment but accelerated dht is disabled
- **providerdelay-datastore** — Same as base provider search delay experiment but the datastore is limited to 1GB to force block retrieval
- **providerdelay-long** — Same as base provider search delay experiment but using much longer delays (10/20/30s)
- **peerset** — Latest kubo comparing settings for number of peer connections





Thunderdome Future



IPFS Camp

Roadmap

- Q4 2022 - Automatic Release Comparisons
 - Test every Kubo release candidate against previous release
- Q1 2023 - Single Click Experiments
 - Allow collaborators to create and run experiments
- Q2 2023 - Alternate Gateway Implementations
 - Experiments with iroh, pomegranate etc.



Long term vision

- More target types - not just IPFS gateway implementations
- More load sources - support protocols other than HTTP such as bitswap
- Bring your own infra - enable deployment on non-AWS infrastructure
- More experiments... more data...



IPFS Camp



Thunderdome Get Involved



IPFS Camp

Got an experiment?

- Find me on Filecoin slack in [#prodeng](#)
- Or on IPFS discord in [#probe-lab](#)
- Raise an issue in the thunderdome Github repo:

<https://github.com/ipfs-shipyard/thunderdome/issues>

(soon you can run your own)



IPFS Camp

Get the source code

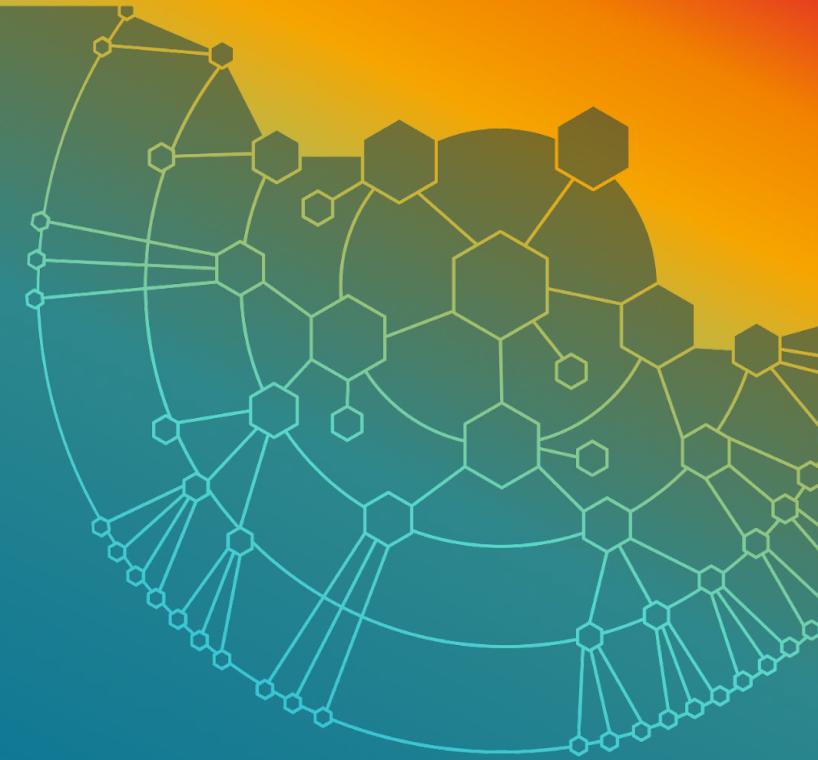
- It's open source on Github

<https://github.com/ipfs-shipyard/thunderdome>



IPFS Camp

Thank you



IPFS Camp